

Foundations of Data Science

Lecture 4: Latent variable models

Prof. Gilles Louppe
g.louppe@uliege.be



What if we had not been given the species labels? Which unobserved quantities would explain what we see?

Probabilistic modeling of data

Recap from Lecture 2

A measurement is what a process produces from an entity $\omega \in \Omega$ (a penguin) under measurement conditions $\xi \in \Xi$ (the scale, the observer, the day),

$$f : \Omega \times \Xi \rightarrow \mathcal{X}.$$

Neither ω nor ξ is observed, only $\mathbf{x} = f(\omega, \xi)$.

Given a distribution $p(\omega, \xi)$ over entities and conditions, f induces the **data distribution**

$$p_r(\mathbf{x}) = \iint p(\omega, \xi) \delta(\mathbf{x} - f(\omega, \xi)) d\omega d\xi.$$

A **parametric probabilistic model** encodes assumptions about how data are generated. It is specified by a parametric family

$$\mathcal{P} = \{p(\mathbf{x} \mid \theta) : \theta \in \Theta\},$$

where $p(\mathbf{x} \mid \theta)$ is a probability distribution over $\mathcal{X}$, θ are parameters, and Θ is the parameter space.

$p(\mathbf{x} \mid \theta)$ is not the data distribution $p_r(\mathbf{x})$, only a model of it! A *true* parameter exists only if p_r belongs to the family $\mathcal{P}$.



Example. Body masses of penguins could be modeled as

$$p(x \mid \mu, \sigma^2) = \mathcal{N}(x \mid \mu, \sigma^2),$$

where the parameters are $\theta = (\mu, \sigma^2)$, with $\mu > 0$ and $\sigma^2 > 0$.

This would assume

- body masses cluster around a central value μ ,
- variability is symmetric and controlled by σ^2 ,
- extreme values are rare (body masses are normally distributed).



Frequentist inference

In the Frequentist framework, θ is an unknown but fixed quantity. The data are assumed to be generated from the model for some parameter value θ^* ,

$$\mathbf{x}_{\text{obs}} \sim p(\mathbf{x} \mid \theta^*).$$

An **estimator** is a recipe $\hat{\theta}(\mathbf{x})$ that turns any data set into a value of θ , for example the maximum likelihood estimator

$$\hat{\theta}(\mathbf{x}) = \arg \max_{\theta} p(\mathbf{x} \mid \theta).$$

Applied to $\mathbf{x}_{\text{obs}}$, it gives a **point estimate**. For the penguins, $\hat{\theta} = (\hat{\mu}, \hat{\sigma}^2)$ with $\hat{\mu} = 4202$ g and $\hat{\sigma} = 801$ g.

Sampling distribution

A point estimate says nothing about its precision. Another colony of 342 penguins would have given another estimate.

Since $\mathbf{x} \sim p(\mathbf{x} \mid \theta^*)$ is random, so is $\hat{\theta}(\mathbf{x})$. Its distribution is the **sampling distribution** of the estimator.

For the sample mean, by the central limit theorem,

$$\hat{\mu}(\mathbf{x}) = \frac{1}{N} \sum_{i=1}^N x_i \sim \mathcal{N} \left(\mu^*, \frac{\sigma^{*2}}{N} \right) \quad \text{approximately.}$$

Its standard deviation $\sigma^* / \sqrt{N}$ is the **standard error**: how much the estimate varies from one data set to the next.

σ^* is unknown too, so it is replaced by its estimate. For the penguins, the estimated standard error is $\hat{\sigma} / \sqrt{N} = 801 / \sqrt{342} \approx 43$ g.

Confidence interval

A 95% **confidence interval** is a set $C(\mathbf{x})$ computed from the data such that, for every θ^* ,

$$P_{\mathbf{x} \sim p(\mathbf{x}|\theta^*)} (\theta^* \in C(\mathbf{x})) = 0.95.$$

The **likelihood ratio** compares each θ with the maximum likelihood estimate,

$$\lambda(\theta; \mathbf{x}) = -2 \log \frac{p(\mathbf{x} | \theta)}{p(\mathbf{x} | \hat{\theta}(\mathbf{x}))}.$$

Wilks' theorem states that if $\mathbf{x} \sim p(\mathbf{x} | \theta^*)$, then $\lambda(\theta^*; \mathbf{x})$ converges in distribution to χ_k^2 as $N \rightarrow \infty$, whatever θ^* .

Wilks' theorem does not require knowing θ^* . With c the 95% quantile of χ_k^2 ,

$$P_{\mathbf{x} \sim p(\mathbf{x}|\theta^*)} (\lambda(\theta^*; \mathbf{x}) \leq c) \approx 0.95 \quad \text{whatever } \theta^*.$$

Define the set of all parameter values that satisfy the same inequality,

$$C(\mathbf{x}) = \{\theta : \lambda(\theta; \mathbf{x}) \leq c\}.$$

Taking $\theta = \theta^*$,

$$\theta^* \in C(\mathbf{x}) \iff \lambda(\theta^*; \mathbf{x}) \leq c.$$

For every data set $\mathbf{x}$, the two statements are both true or both false. They are the same event, with the same probability,

$$P_{\mathbf{x} \sim p(\mathbf{x}|\theta^*)} (\theta^* \in C(\mathbf{x})) = P_{\mathbf{x} \sim p(\mathbf{x}|\theta^*)} (\lambda(\theta^*; \mathbf{x}) \leq c) \approx 0.95.$$

Therefore $C(\mathbf{x})$ is a 95% confidence interval, approximately for large N .

Example: mean body mass


Let $x_i \sim \mathcal{N}(\mu, \sigma^2)$, with $\theta = \mu$ (so $k = 1$) and σ fixed to $\hat{\sigma} = 801$ g. The maximum likelihood estimate is $\hat{\mu} = \frac{1}{N} \sum_i x_i$, and the likelihood ratio is

$$\lambda(\mu; \mathbf{x}) = -2 \log \frac{p(\mathbf{x} \mid \mu)}{p(\mathbf{x} \mid \hat{\mu})} = \frac{N(\hat{\mu} - \mu)^2}{\sigma^2}.$$

If $\mathbf{x} \sim p(\mathbf{x} \mid \mu^*)$, then $\hat{\mu} \sim \mathcal{N}(\mu^*, \sigma^2/N)$, and

$$\lambda(\mu^*; \mathbf{x}) = \left(\frac{\hat{\mu} - \mu^*}{\sigma/\sqrt{N}} \right)^2$$

is the square of a standard Gaussian. Wilks' theorem is exact here: $\lambda(\mu^*; \mathbf{x})$ is χ_1^2 distributed for every N , and its 95% quantile is $c = 1.96^2 = 3.84$.



The confidence interval is

$$C(\mathbf{x}) = \{\mu : \lambda(\mu; \mathbf{x}) \leq 3.84\} = \hat{\mu} \pm 1.96 \frac{\sigma}{\sqrt{N}} = [4117, 4287] \text{ g.}$$




50 colonies simulated from the model with $\mu^* = 4200 \text{ g}$. Three intervals miss μ^* .

95% of the intervals built by C contain μ^* . Whether $[4117, 4287] \text{ g}$ does is unknown.

Wilks' theorem can be replaced by simulation. For any candidate θ , the forward model gives the distribution of $\lambda(\theta; \mathbf{x})$ when the data come from $p(\mathbf{x} \mid \theta)$:

1. simulate data sets $\mathbf{x}_1, \dots, \mathbf{x}_M \sim p(\mathbf{x} \mid \theta)$,
2. compute $\lambda(\theta; \mathbf{x}_m)$ for each of them,
3. denote their empirical 95% quantile as $c(\theta)$.

At $\theta = \theta^*$, this is the distribution of $\lambda(\theta^*; \mathbf{x})$, and $c(\theta^*)$ takes the place of c .




One candidate, $\mu = 4150$ g: $c(\mu) = 3.80$, against 3.84 for χ_1^2 .

The construction is then the same as before, with $c(\theta)$ in place of c ,

$$C(\mathbf{x}) = \{\theta : \lambda(\theta; \mathbf{x}) \leq c(\theta)\}.$$

It contains θ^* exactly when $\lambda(\theta^*; \mathbf{x}) \leq c(\theta^*)$, which happens with probability 0.95, for any N .

On the observed data, we compute $\lambda(\theta; \mathbf{x}_{\text{obs}})$ at every candidate and keep those below their own $c(\theta)$:



$$C(\mathbf{x}_{\text{obs}}) = [4117, 4289] \text{ g by simulation, } [4117, 4287] \text{ g with } \chi_1^2.$$

Bayesian inference

In the Bayesian framework, θ is treated as a random variable with prior distribution $p(\theta)$ encoding beliefs about plausible parameter values before observing any data.

A **Bayesian model** therefore specifies a joint distribution over data and parameters,

$$p(\mathbf{x}, \theta) = p(\mathbf{x} \mid \theta)p(\theta),$$


where

- $p(\mathbf{x} \mid \theta)$ is the likelihood,
- $p(\theta)$ is the prior over parameters.

The prior $p(\theta)$ is often harder to justify than the forward model $p(\mathbf{x} \mid \theta)$. A **prior predictive check** simulates data from the model before any data are observed,

$$\mathbf{x} \sim p(\mathbf{x}) = \int p(\mathbf{x} \mid \theta) p(\theta) d\theta,$$

and compares them with what is known about the data.



Five colonies of 342 penguins simulated from the prior (top), and from a wider prior on σ^2 (bottom), against the penguins actually measured.

Fitting a Bayesian model to observed data $\mathbf{x}_{\text{obs}}$ consists in computing the posterior distribution of the parameters given the data. Using Bayes' rule,

$$p(\theta \mid \mathbf{x}_{\text{obs}}) = \frac{p(\mathbf{x}_{\text{obs}} \mid \theta)p(\theta)}{p(\mathbf{x}_{\text{obs}})}.$$

Depending on the structure of the model, this computation may be easy, difficult, or even intractable.

A 95% **credible interval** is a set C such that

$$P(\theta \in C \mid \mathbf{x}_{\text{obs}}) = 0.95.$$

The probability is over θ . The data are fixed: they are the ones observed.

Confidence or credible?

For the penguins, take σ fixed to $\hat{\sigma} = 801$ g and a uniform prior on μ between 0 and 10 kg. The posterior is then

$$p(\mu \mid \mathbf{x}_{\text{obs}}) \propto \mathcal{N}\left(\mu \mid \hat{\mu}, \frac{\sigma^2}{N}\right) 1[0 \leq \mu \leq 10000],$$

and its 95% credible interval is $[4117, 4287]$ g, the confidence interval again. The two intervals differ in meaning:

Confidence interval

- the data are random, μ^* is fixed
- 95% of the intervals built this way contain μ^*
- needs a model

Credible interval

- the data are fixed, μ is random
- μ is believed to lie in $[4117, 4287]$ g with probability 0.95
- needs a model and a prior

Latent variable models

Models of data generating processes are often written as a cascade of variables, each produced from the ones before it. Only the last ones are observed, through a measurement process; the others are latent:

$$\underbrace{\theta \rightarrow \mathbf{z}_1 \rightarrow \cdots \rightarrow \mathbf{z}_L}_{\text{latent}} \rightarrow \underbrace{\mathbf{x}}_{\text{measured}}$$

- **Particle physics**: theory $\rightarrow$ collision $\rightarrow$ particle showers $\rightarrow$ detector hits.
- **Epidemiology**: transmission $\rightarrow$ infections $\rightarrow$ symptoms $\rightarrow$ tests $\rightarrow$ cases.
- **Hydrology**: rainfall $\rightarrow$ soil saturation $\rightarrow$ Meuse discharge $\rightarrow$ water level.
- **Astronomy**: stellar mass $\rightarrow$ luminosity $\rightarrow$ brightness at Earth $\rightarrow$ pixels.

Joint distribution

A **latent variable model** is a probabilistic model that assumes unobserved (latent) variables $\mathbf{z}$ that mediate the relationship between observed data $\mathbf{x}$ and model parameters θ .

It specifies a joint distribution over observed variables, latent variables, and parameters,

$$p(\mathbf{x}, \mathbf{z}, \theta) = p(\mathbf{x} \mid \mathbf{z}, \theta)p(\mathbf{z} \mid \theta)p(\theta),$$

where $\mathbf{x}$ is the observed data, $\mathbf{z}$ are the latent variables, and θ are the parameters.

More generally, for a dataset of N observations $\{\mathbf{x}_1, \dots, \mathbf{x}_N\}$, a latent variable model specifies a joint distribution

$$p(\mathbf{x}_{1:N}, \mathbf{z}_{1:N}, \theta) = \left(\prod_{i=1}^N p(\mathbf{x}_i \mid \mathbf{z}_i, \theta) p(\mathbf{z}_i \mid \theta) \right) p(\theta),$$

where $\mathbf{z}_i$ are the latent variables associated with observation $\mathbf{x}_i$.

The factorization assumes that each observation $\mathbf{x}_i$ is mediated by its own latent variable $\mathbf{z}_i$, and that observations are conditionally independent given their latent variables and the parameters. All are governed by shared parameters θ .

Marginal likelihood

The latent variables are not observed. Integrating them out gives the likelihood of the data alone,

$$p(\mathbf{x} \mid \theta) = \int p(\mathbf{x} \mid \mathbf{z}, \theta) p(\mathbf{z} \mid \theta) d\mathbf{z},$$

or, for N conditionally independent observations,

$$p(\mathbf{x}_{1:N} \mid \theta) = \prod_{i=1}^N \int p(\mathbf{x}_i \mid \mathbf{z}_i, \theta) p(\mathbf{z}_i \mid \theta) d\mathbf{z}_i.$$

This integral is where the difficulty of latent variable models lies: it is what maximum likelihood maximizes, and it is the normalizer of the posterior over the latent variables.


Graphical model representation

A latent variable model can be drawn as a **graphical model**: one node per variable (observed, latent, or parameter), and an arrow $a \rightarrow b$ when the distribution of b depends on a .

The joint distribution is the product, over all variables, of the distribution of each variable given the variables pointing to it.

Read along the arrows, the graph is the generative process. A missing arrow is an assumption.

A latent variable model with $N = 3$ observations, as a graphical model:



$$p(\mathbf{x}_{1:3}, \mathbf{z}_{1:3}, \theta) = \left(\prod_{i=1}^3 p(\mathbf{x}_i \mid \mathbf{z}_i) p(\mathbf{z}_i \mid \theta) \right) p(\theta)$$

Shaded nodes represent observed variables, unshaded nodes represent latent variables or parameters.

Here $\mathbf{x}_i$ depends on θ only through $\mathbf{z}_i$. In general, θ may also point directly at $\mathbf{x}_i$.




Plate notation can be used to compactly represent repeated structures in the graphical model.

Hyperparameters

Conditional distributions in a latent variable model may depend on additional parameters called **hyperparameters**, denoted $\alpha, \beta, \dots$. These are assumed to be fixed nonrandom quantities¹.


For instance, the prior distribution of parameters may depend on hyperparameters,

$$p(\theta \mid \alpha),$$

or the prior distribution of latent variables may depend on hyperparameters,

$$p(\mathbf{z} \mid \theta, \beta).$$

¹: Estimating hyperparameters from data is possible; this is empirical Bayes, discussed in Lecture 8.



Small squares denote fixed hyperparameters.

Inference

Fitting a latent variable model to observed data $\mathbf{x}_{\text{obs}}$ consists in computing the posterior distribution of latent variables and parameters given the data. Using Bayes' rule,


$$p(\mathbf{z}, \theta \mid \mathbf{x}_{\text{obs}}) = \frac{p(\mathbf{x}_{\text{obs}} \mid \mathbf{z}, \theta)p(\mathbf{z} \mid \theta)p(\theta)}{p(\mathbf{x}_{\text{obs}})}.$$

The posterior distribution is used to examine the particular hidden structure that is manifested in the observed data. It can also be used to make predictions about new, unseen data, through the posterior predictive distribution,

$$p(\mathbf{x}_{\text{new}} \mid \mathbf{x}_{\text{obs}}) = \iint p(\mathbf{x}_{\text{new}} \mid \mathbf{z}, \theta)p(\mathbf{z}, \theta \mid \mathbf{x}_{\text{obs}}) d\mathbf{z} d\theta.$$

Example 1: (Probabilistic) PCA

In probabilistic PCA, each observation $\mathbf{x}_i \in \mathbb{R}^d$ is assumed to be generated from a lower-dimensional latent variable $\mathbf{z}_i \in \mathbb{R}^m$ through a linear transformation plus Gaussian noise.



The joint distribution $p(\mathbf{z}, \mathbf{x} \mid \mathbf{B}, \boldsymbol{\mu}, \sigma^2)$ factorizes as $p(\mathbf{z})p(\mathbf{x} \mid \mathbf{z}, \mathbf{B}, \boldsymbol{\mu}, \sigma^2)$, where

- $p(\mathbf{z}) = \mathcal{N}(\mathbf{z} \mid \mathbf{0}, \mathbf{I})$ assumes latent variables are standard Gaussian,
- $p(\mathbf{x} \mid \mathbf{z}, \mathbf{B}, \boldsymbol{\mu}, \sigma^2) = \mathcal{N}(\mathbf{x} \mid \mathbf{B}\mathbf{z} + \boldsymbol{\mu}, \sigma^2\mathbf{I})$ assumes a linear Gaussian observation model, with $\mathbf{B} \in \mathbb{R}^{d \times m}$ the loading matrix, $\boldsymbol{\mu} \in \mathbb{R}^d$ the mean vector, and σ^2 the noise variance.

Therefore, using Gaussian identities, the joint distribution is Gaussian and can be written as

$$p(\mathbf{z}, \mathbf{x} \mid \mathbf{B}, \boldsymbol{\mu}, \sigma^2) = \mathcal{N} \left(\begin{bmatrix} \mathbf{z} \\ \mathbf{x} \end{bmatrix} \mid \begin{bmatrix} \mathbf{0} \\ \boldsymbol{\mu} \end{bmatrix}, \begin{bmatrix} \mathbf{I} & \mathbf{B}^T \\ \mathbf{B} & \mathbf{B}\mathbf{B}^T + \sigma^2\mathbf{I} \end{bmatrix} \right).$$

The posterior distribution $p(\mathbf{z} \mid \mathbf{x}, \mathbf{B}, \boldsymbol{\mu}, \sigma^2)$ is also Gaussian,

$$p(\mathbf{z} \mid \mathbf{x}, \mathbf{B}, \boldsymbol{\mu}, \sigma^2) = \mathcal{N}(\mathbf{z} \mid \boldsymbol{\mu}_{z|x}, \boldsymbol{\Sigma}_{z|x}),$$

where

- $\boldsymbol{\mu}_{z|x} = \mathbf{B}^T(\mathbf{B}\mathbf{B}^T + \sigma^2\mathbf{I})^{-1}(\mathbf{x} - \boldsymbol{\mu})$ is the posterior mean,
- $\boldsymbol{\Sigma}_{z|x} = \mathbf{I} - \mathbf{B}^T(\mathbf{B}\mathbf{B}^T + \sigma^2\mathbf{I})^{-1}\mathbf{B}$ is the posterior covariance.

When $\sigma^2 \rightarrow 0$, $\mathbf{B}^T(\mathbf{B}\mathbf{B}^T + \sigma^2\mathbf{I})^{-1} \rightarrow \mathbf{B}^+$, the pseudo-inverse of $\mathbf{B}$, so

- $\boldsymbol{\mu}_{z|x} \rightarrow \mathbf{B}^+(\mathbf{x} - \boldsymbol{\mu})$. If the columns of $\mathbf{B}$ are orthonormal, $\mathbf{B}^+ = \mathbf{B}^T$ and this is the PCA projection of $\mathbf{x}$.
- $\boldsymbol{\Sigma}_{z|x} \rightarrow \mathbf{I} - \mathbf{B}^+\mathbf{B} = \mathbf{0}$: the posterior collapses to a point mass at that projection.

Probabilistic PCA recovers classical PCA in the limit of vanishing noise!

The parameters $\theta = (\mathbf{B}, \boldsymbol{\mu}, \sigma^2)$ can be estimated from data $\mathbf{x}_{1:N}$ by maximum (marginal) likelihood,

$$(\hat{\mathbf{B}}, \hat{\boldsymbol{\mu}}, \hat{\sigma}^2) = \arg \max_{\mathbf{B}, \boldsymbol{\mu}, \sigma^2} \prod_{i=1}^N p(\mathbf{x}_i \mid \mathbf{B}, \boldsymbol{\mu}, \sigma^2),$$

where $p(\mathbf{x} \mid \mathbf{B}, \boldsymbol{\mu}, \sigma^2) = \int p(\mathbf{x} \mid \mathbf{z}, \mathbf{B}, \boldsymbol{\mu}, \sigma^2) p(\mathbf{z}) d\mathbf{z}$ is the marginal likelihood.

Since the joint distribution is Gaussian, the marginal likelihood is also Gaussian,

$$p(\mathbf{x} \mid \mathbf{B}, \boldsymbol{\mu}, \sigma^2) = \mathcal{N}(\mathbf{x} \mid \boldsymbol{\mu}, \mathbf{B}\mathbf{B}^T + \sigma^2\mathbf{I}).$$

Therefore, writing $\Sigma = \mathbf{B}\mathbf{B}^T + \sigma^2\mathbf{I}$, maximum likelihood estimation reduces to

$$(\hat{\mathbf{B}}, \hat{\boldsymbol{\mu}}, \hat{\sigma}^2) = \arg \max_{\mathbf{B}, \boldsymbol{\mu}, \sigma^2} \prod_{i=1}^N \mathcal{N}(\mathbf{x}_i \mid \boldsymbol{\mu}, \Sigma).$$

The solution can be derived in closed form, yielding


- $\hat{\boldsymbol{\mu}} = \frac{1}{N} \sum_{i=1}^N \mathbf{x}_i$ (the sample mean),
- $\hat{\mathbf{B}} = \mathbf{U}_m(\boldsymbol{\Lambda}_m - \hat{\sigma}^2\mathbf{I})^{1/2}\mathbf{R}$, where $\mathbf{U}_m$ holds the top m eigenvectors of the sample covariance matrix $\mathbf{S} = \frac{1}{N} \sum_{i=1}^N (\mathbf{x}_i - \hat{\boldsymbol{\mu}})(\mathbf{x}_i - \hat{\boldsymbol{\mu}})^T$, $\boldsymbol{\Lambda}_m$ the corresponding eigenvalues, and $\mathbf{R}$ is an arbitrary rotation matrix,
- $\hat{\sigma}^2 = \frac{1}{d-m} \sum_{j=m+1}^d \lambda_j$, where λ_j are the eigenvalues of $\mathbf{S}$.

Since $\mathbf{R}$ is free, $\mathbf{B}$ is identified only up to a rotation: the model determines the latent subspace, and any orthonormal basis of it fits the data equally well.



Example 2: Mixture models

Mixture models assume that data are generated from a mixture of several underlying distributions, each corresponding to a different cluster or component.



For a Gaussian mixture model with K components, each observation $\mathbf{x}_i \in \mathbb{R}^d$ is assumed to be generated by first selecting a component $z_i \in \{1, \dots, K\}$ according to a categorical distribution, then sampling $\mathbf{x}_i$ from a Gaussian distribution associated with that component.

The parameters are $\theta = (\boldsymbol{\pi}, \boldsymbol{\mu}_{1:K}, \sigma_{1:K}^2)$, where $\boldsymbol{\pi}$ holds the mixture weights. The joint distribution $p(\theta, z_{1:N}, \mathbf{x}_{1:N} \mid \alpha, \sigma_\mu^2, \sigma_\sigma^2)$ factorizes as

$$p(\boldsymbol{\pi} \mid \alpha) \prod_{k=1}^K p(\boldsymbol{\mu}_k \mid \sigma_\mu^2) p(\sigma_k^2 \mid \sigma_\sigma^2) \prod_{i=1}^N p(z_i \mid \boldsymbol{\pi}) p(\mathbf{x}_i \mid z_i, \boldsymbol{\mu}_{z_i}, \sigma_{z_i}^2),$$

where


- $p(\boldsymbol{\pi} \mid \alpha) = \text{Dirichlet}(\boldsymbol{\pi} \mid \alpha)$ is the prior over mixture weights,
- $p(\boldsymbol{\mu}_k \mid \sigma_\mu^2) = \mathcal{N}(\boldsymbol{\mu}_k \mid \mathbf{0}, \sigma_\mu^2 \mathbf{I})$ is the prior over component means,
- $p(\sigma_k^2 \mid \sigma_\sigma^2) = \text{Lognormal}(\sigma_k^2 \mid 0, \sigma_\sigma^2)$ is the prior over component variances,
- $p(z_i \mid \boldsymbol{\pi}) = \text{Categorical}(z_i \mid \boldsymbol{\pi})$ is the distribution over components,
- $p(\mathbf{x}_i \mid z_i, \boldsymbol{\mu}_{z_i}, \sigma_{z_i}^2) = \mathcal{N}(\mathbf{x}_i \mid \boldsymbol{\mu}_{z_i}, \sigma_{z_i}^2 \mathbf{I})$ is the Gaussian observation model.

Computing the posterior $p(\theta, z_{1:N} \mid \mathbf{x}_{1:N}, \alpha, \sigma_\mu^2, \sigma_\sigma^2)$ amounts to clustering: each component is a cluster, and z_i is the cluster of observation i .

This posterior has no closed form. Its normalizing constant

$$p(\mathbf{x}_{1:N}) = \int \sum_{z_1=1}^K \cdots \sum_{z_N=1}^K p(\theta, z_{1:N}, \mathbf{x}_{1:N}) d\theta$$

sums over all K^N assignments of the observations to components: for 342 penguins and 3 components, about 10^{163} terms.




Three Gaussians fitted to the bill measurements by maximum likelihood.
Each penguin is coloured by its most likely component.

Example 3: Mixed membership models

Nested sets of latent variables can also be used to model more complex generative structures.

For instance, in mixed membership models of text documents (**latent Dirichlet allocation**), each document is assumed to be generated from a mixture of topics, where each topic is characterized by a distribution over words.



Posterior inference in mixed membership models can be used to discover the underlying topics in a corpus of documents and to infer the topic proportions for each document.

gene 0.04
dna 0.02
genetic 0.01
...

life 0.02
evolve 0.01
organism 0.01
...

brain 0.04
neuron 0.02
nerve 0.01
...


data 0.02
number 0.02
computer 0.01
...

Seeking Life's Bare (Genetic) Necessities

COLD SPRING HARBOR, NEW YORK—How many **genes** does an **organism** need to **survive**? Last week at the genome meeting here,* two genome researchers with radically different approaches presented complementary views of the basic genes needed for **life**. One research team, using **computer** analyses to compare known **genomes**, concluded that today's **organisms** can be sustained with just 250 genes, and that the earliest life forms required a mere 128 **genes**. The other researcher mapped genes in a simple parasite and estimated that for this organism, 800 genes are plenty to do the job—but that anything short of 100 wouldn't be enough.

Although the numbers don't match precisely, those **predictions**

"are not all that far apart," especially in comparison to the 75,000 **genes** in the human genome, notes Siv Andersson at Uppsala University in Sweden, who arrived at the 800 number. But coming up with a consensus answer may be more than just a **genetic** numbers game, particularly as more and more **genomes** are completely mapped and sequenced. "It may be a way of organizing any newly **sequenced genome**," explains Arcady Mushegian, a **computational** molecular biologist at the National Center for Biotechnology Information (NCBI) in Bethesda, Maryland. Comparing an



* Genome Mapping and Sequencing, Cold Spring Harbor, New York, May 8 to 12.

Stripping down. Computer analysis yields an estimate of the minimum modern and ancient genomes.

SCIENCE • VOL. 272 • 24 MAY 1996

Figure 1: **The intuitions behind latent Dirichlet allocation.** We assume that some number of “topics,” which are distributions over words, exist for the whole collection (far left). Each document is assumed to be generated as follows. First choose a distribution over the topics (the histogram at right); then, for each word, choose a topic assignment (the colored coins) and choose the word from the corresponding topic. The topics and topic assignments in this figure are illustrative—they are not fit from real data. See Figure 2 for topics fit from data.





Figure 2: **Real inference with LDA.** We fit a 100-topic LDA model to 17,000 articles from the journal *Science*. At left is the inferred topic proportions for the example article in Figure 1. At right are the top 15 most frequent words from the most frequent topics found in this article.

Example 4: How far are the stars?

Gaia measured the **parallax** of 1.47 billion stars: the angle that the radius of Earth's orbit subtends at the star, which shrinks as the star gets further away,


$$\varpi = \frac{1}{r},$$

with ϖ in milliarcseconds (a thousandth of an arcsecond) and r in kiloparsecs (about 3260 light-years).



Distances are what turn a catalogue of angles and brightnesses into physics:

- how bright a star looks depends on how bright it really is and on how far away it is; the distance separates the two, and how bright it really is tells its mass and its age,
- how fast it drifts across the sky, an angle per year, becomes a speed in kilometres per second, which is how the mass of the Galaxy, dark matter included, is weighed,
- the catalogue stops being a picture of the sky and becomes a three-dimensional map.



5000 stars from Gaia DR3, each parallax ϖ_i reported with an uncertainty σ_i .

For 82% of them $\varpi_i < 5\sigma_i$; for the 17% below the line, $1/\varpi_i$ is negative.

Geometry. The parallax of a star at distance r_i is exactly $1/r_i$. No modelling freedom here: it is what a parallax is.

The instrument. Gaia reports an estimate of that angle, with an uncertainty σ_i computed star by star. Modelling its errors as unbiased and Gaussian gives


$$p(\varpi_i \mid r_i, \sigma_i) = \mathcal{N}(\varpi_i \mid 1/r_i, \sigma_i^2),$$

which is why a measured parallax can be negative while a distance cannot.

The Galaxy. The number of stars at distance r_i grows with the room available and decreases with their density:

$$p(r_i \mid L) \propto \underbrace{r_i^2}_{\text{room at distance } r_i} \times \underbrace{\exp(-r_i/L)}_{\text{density of stars}}, \quad r_i > 0.$$

The scale length L of the decay is unknown too, and shared by all stars, so it gets a vague prior of its own, $p(L) = \text{Uniform}(L \mid 0, 5)$.



Everything unobserved is inferred at once,

$$p(r_{1:N}, L \mid \varpi_{1:N}, \sigma_{1:N}) \propto p(L) \prod_{i=1}^N p(\varpi_i \mid r_i, \sigma_i) p(r_i \mid L),$$


a posterior over $N + 1$ unknowns.

What do the stars say about the Galaxy? Integrating out every distance leaves the length scale alone,

$$p(L \mid \varpi_{1:N}, \sigma_{1:N}) \propto p(L) \prod_{i=1}^N \int p(\varpi_i \mid r_i, \sigma_i) p(r_i \mid L) dr_i,$$

the marginal likelihood of the catalogue, times the prior.

Each star contributes one factor. One noisy parallax barely constrains L ; the product over 5000 stars pins it down.




The stellar density thins out along these lines of sight with a scale length of **1.03 ± 0.014** kpc: a measurement of the shape of the Galaxy.

And where is each star? The distance of star i comes from the same posterior, with L and the other distances integrated out,

$$p(r_i \mid \varpi_{1:N}, \sigma_{1:N}) \propto \int p(\varpi_i \mid r_i, \sigma_i) p(r_i \mid L) p(L \mid \varpi_{1:N}, \sigma_{1:N}) dL.$$

When that parallax is precise, it decides the answer on its own. When it is not, the distance comes mostly from what the whole catalogue has established about where stars sit.



Three stars of the sample: how likely each distance is before the parallax is used (grey), and after (blue).

Cheat sheet for Gaussian models (Särkkä, 2013)

If $\mathbf{x}$ and $\mathbf{y}$ have the joint Gaussian distribution

$$p \begin{pmatrix} \mathbf{x} \\ \mathbf{y} \end{pmatrix} = \mathcal{N} \left(\begin{pmatrix} \mathbf{x} \\ \mathbf{y} \end{pmatrix} \middle| \begin{pmatrix} \mathbf{a} \\ \mathbf{b} \end{pmatrix}, \begin{pmatrix} \mathbf{A} & \mathbf{C} \\ \mathbf{C}^T & \mathbf{B} \end{pmatrix} \right),$$

then the marginal and conditional distributions of $\mathbf{x}$ and $\mathbf{y}$ are given by

$$p(\mathbf{x}) = \mathcal{N}(\mathbf{x}|\mathbf{a}, \mathbf{A})$$

$$p(\mathbf{y}) = \mathcal{N}(\mathbf{y}|\mathbf{b}, \mathbf{B})$$

$$p(\mathbf{x}|\mathbf{y}) = \mathcal{N}(\mathbf{x}|\mathbf{a} + \mathbf{CB}^{-1}(\mathbf{y} - \mathbf{b}), \mathbf{A} - \mathbf{CB}^{-1}\mathbf{C}^T)$$

$$p(\mathbf{y}|\mathbf{x}) = \mathcal{N}(\mathbf{y}|\mathbf{b} + \mathbf{C}^T\mathbf{A}^{-1}(\mathbf{x} - \mathbf{a}), \mathbf{B} - \mathbf{C}^T\mathbf{A}^{-1}\mathbf{C}).$$

If the random variables $\mathbf{x}$ and $\mathbf{y}$ have Gaussian probability distributions

$$p(\mathbf{x}) = \mathcal{N}(\mathbf{x}|\mathbf{m}, \mathbf{P})$$
$$p(\mathbf{y}|\mathbf{x}) = \mathcal{N}(\mathbf{y}|\mathbf{H}\mathbf{x} + \mathbf{u}, \mathbf{R}),$$

then the joint distribution of $\mathbf{x}$ and $\mathbf{y}$ is Gaussian with

$$p \begin{pmatrix} \mathbf{x} \\ \mathbf{y} \end{pmatrix} = \mathcal{N} \left(\begin{pmatrix} \mathbf{x} \\ \mathbf{y} \end{pmatrix} \middle| \begin{pmatrix} \mathbf{m} \\ \mathbf{H}\mathbf{m} + \mathbf{u} \end{pmatrix}, \begin{pmatrix} \mathbf{P} & \mathbf{P}\mathbf{H}^T \\ \mathbf{H}\mathbf{P} & \mathbf{H}\mathbf{P}\mathbf{H}^T + \mathbf{R} \end{pmatrix} \right).$$